

Machine Learning for Cognitive Sciences: Principles and Applications

Simona Cocco (Physics Departement of ENS)
simona.cocco@phys.ens.fr CM &TD

Vito Dichio (Physics Departement of ENS)
vito.dichio@phys.ens.fr TD

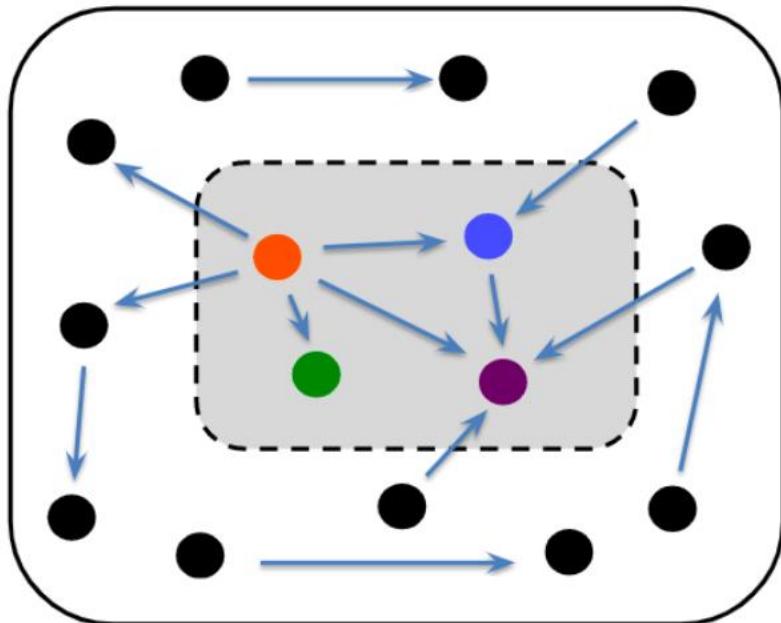
Director of the Cogmaster: Jean Baptiste Andre
jean-baptiste.andre@ens.fr

<https://moodle.psl.eu/enrol/index.php?id=34555>

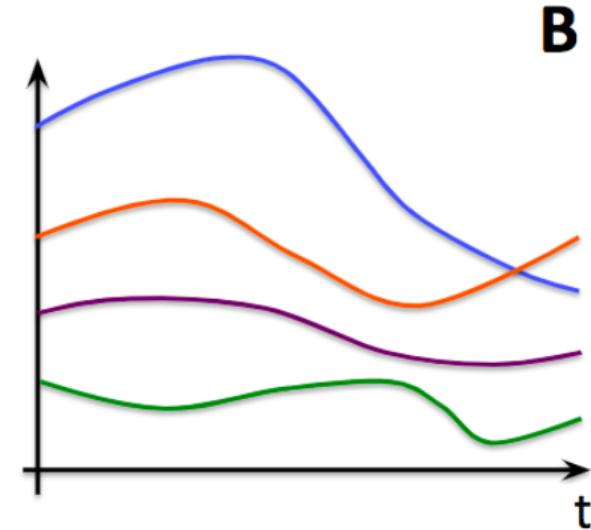
Accès anonyme pwd:MLCS2025

Class: Wednesday salle RIBOT,
Tutorial Thursday salle Langevin

Machine Learning for Cognitive Sciences: Principles and Applications

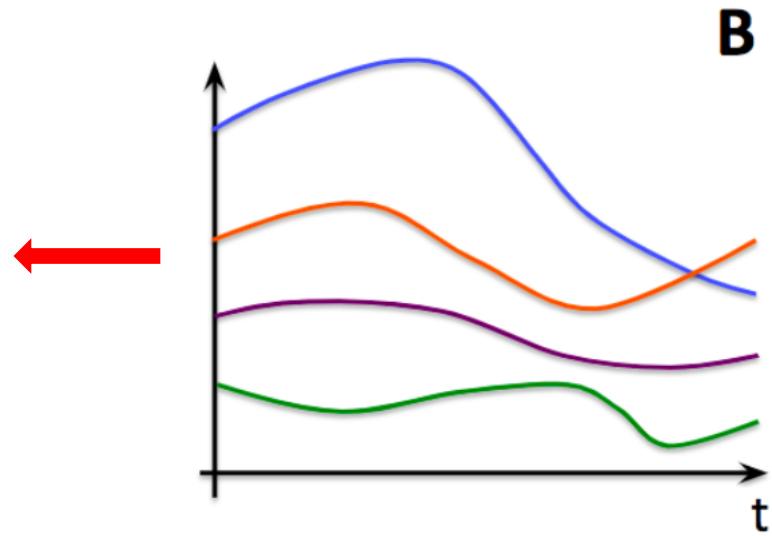
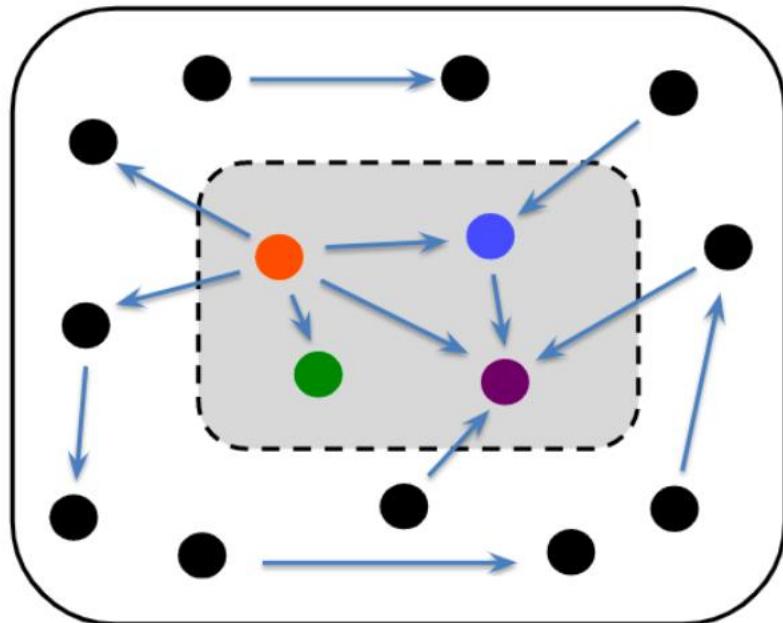


A complex system made of many particles which interact together



Observation of a limited part of the system and of the dynamics of its components

Machine Learning for Cognitive Sciences: Principles and Applications



A complex system made of many particles which interact together

Observation of a limited part of the system and of the dynamics of its components

Aim: Understand how the system works, Predict the future behavior of it.

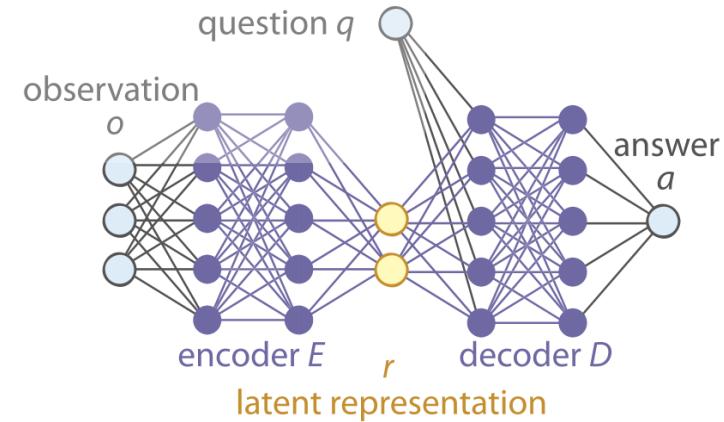
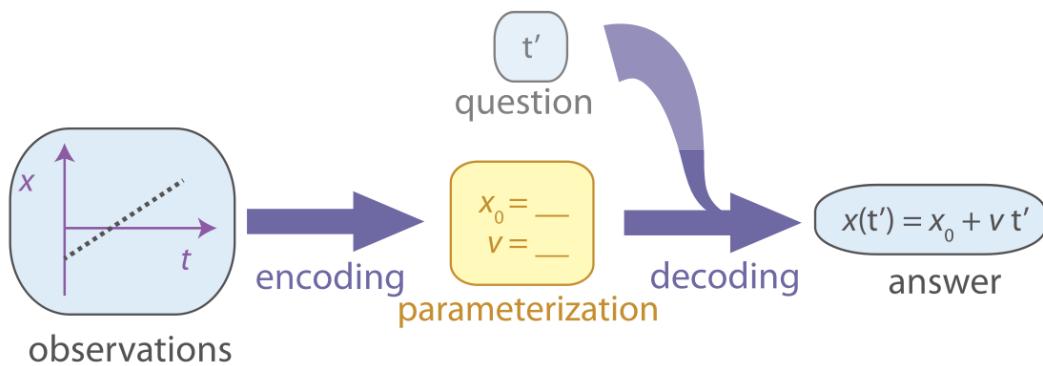
Machine Learning for Cognitive Sciences:

Principles and Applications

Exemples of ML utilization in Neuroscience (Neural recordings), Images (MNIST) and Language Classification (Newspaper texts), behavioral Data (trajectories of zebrafish).

ML as a learning and coding paradigm for real neural networks .

Machine Learning for Cognitive Sciences: Principles and Applications



Auto-Encoder with SciNet

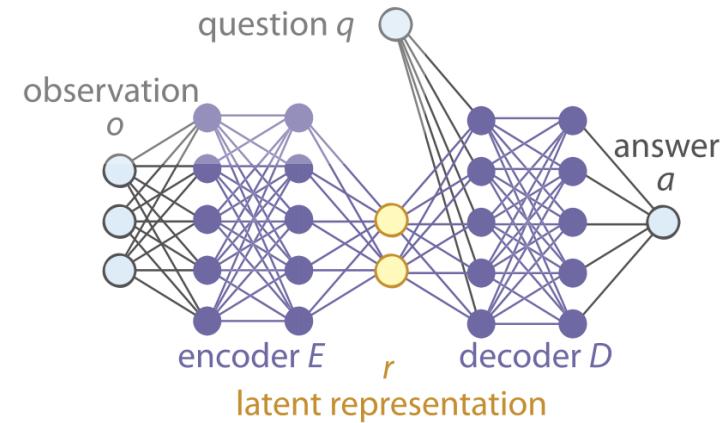
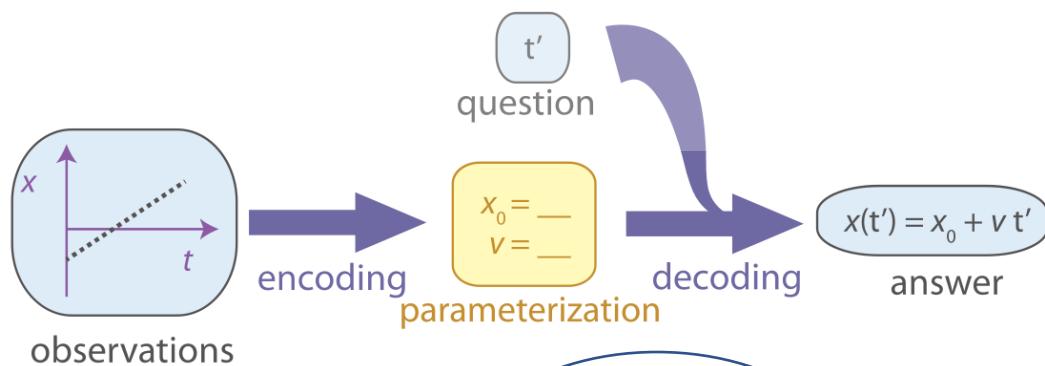
Use the network to

predict the behavior of the system

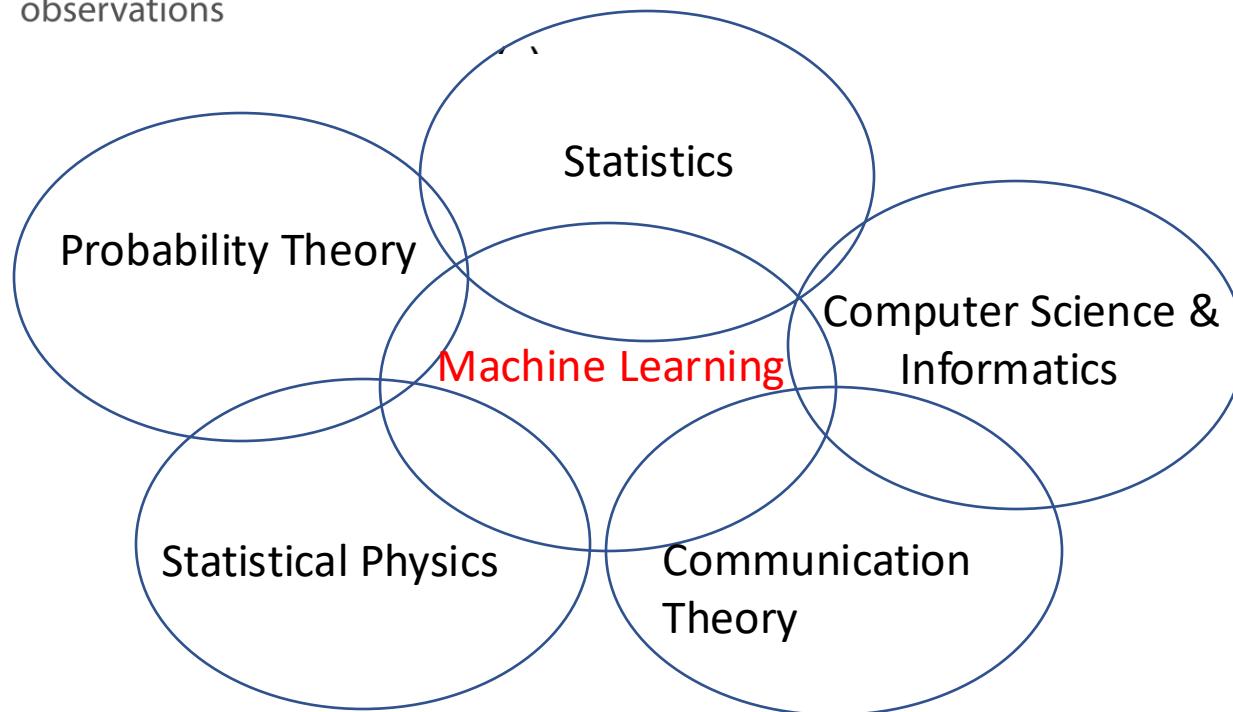
Interpret the parameters: how they encode the input

Find the best network: related to the number of data

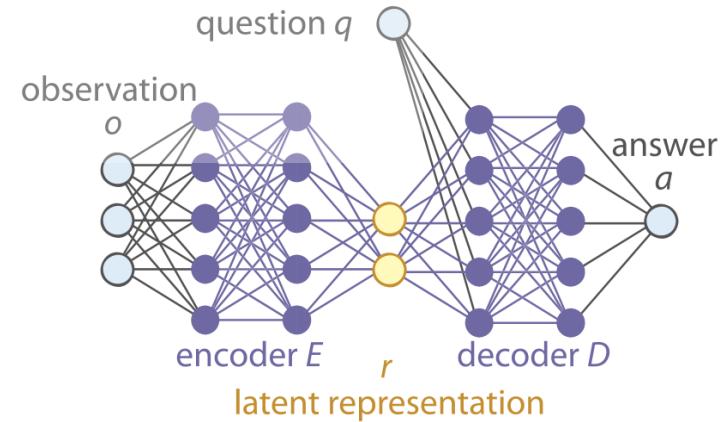
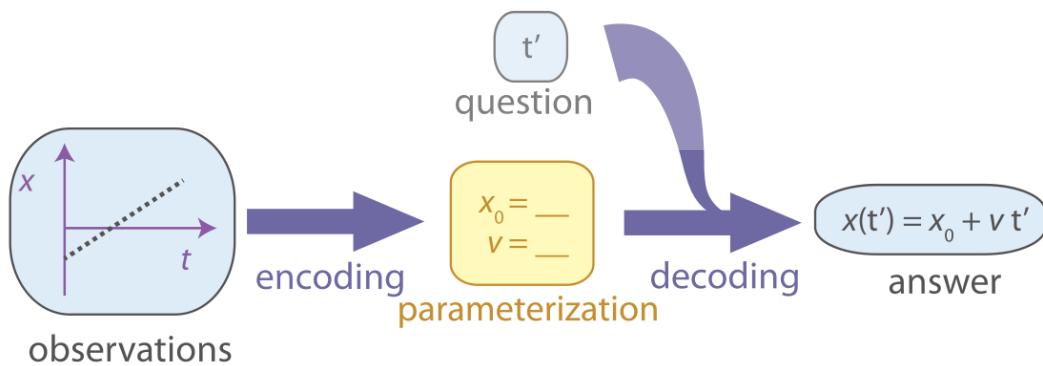
Machine Learning for Cognitive Sciences: Principles and Applications



Auto-Encoder with SciNet



Machine Learning for Cognitive Sciences: Principles and Applications



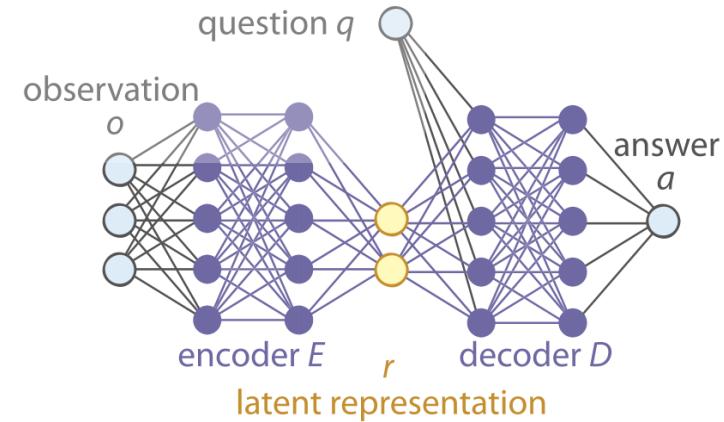
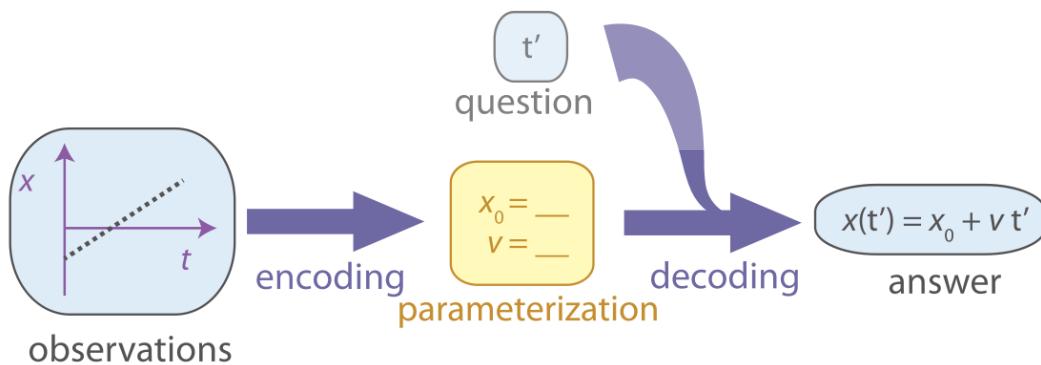
Auto-Encoder with SciNet

Aims of the Course:

- Learn **Bayesian framework** for the Inference.
- Learn Inference Algorithm, **Programming** yourself in Python to avoid black box effect,
- Control learning: **Statistical Physics & Statistics** to understand how to adapt the inferred models to the quality/quantity of data: eg. avoid **Overfitting** (use **Regularizations**) , Compare with **Null models**, Use **Cross Validations**.

Machine Learning for Cognitive Sciences:

Principles and Applications



Auto-Encoder with SciNet

Aims of the Course:

- Learn **Bayesian framework** for the Inference.
- Learn Inference Algorithm, **Programming** yourself in Python to avoid black box effect,
- Control learning: **Statistical Physics & Statistics** to understand how to adapt the inferred models to the quality/quantity of data: eg. avoid **Overfitting** (use **Regularizations**) , Compare with **Null models**, Use **Cross Validations**.
- **Unsupervised Learning: Graphical models, Multi-Layer networks .**
- **Clustering and low dimensional representation of data** (eg. Principal Component Analysis)
- **Supervised Learning:** Neural Network for classification (**Perceptron algorithm**).
- **Time-dependent process:** Markov-Model and Hidden Markov-Model.

Machine Learning for Cognitive Sciences: Principles and Applications

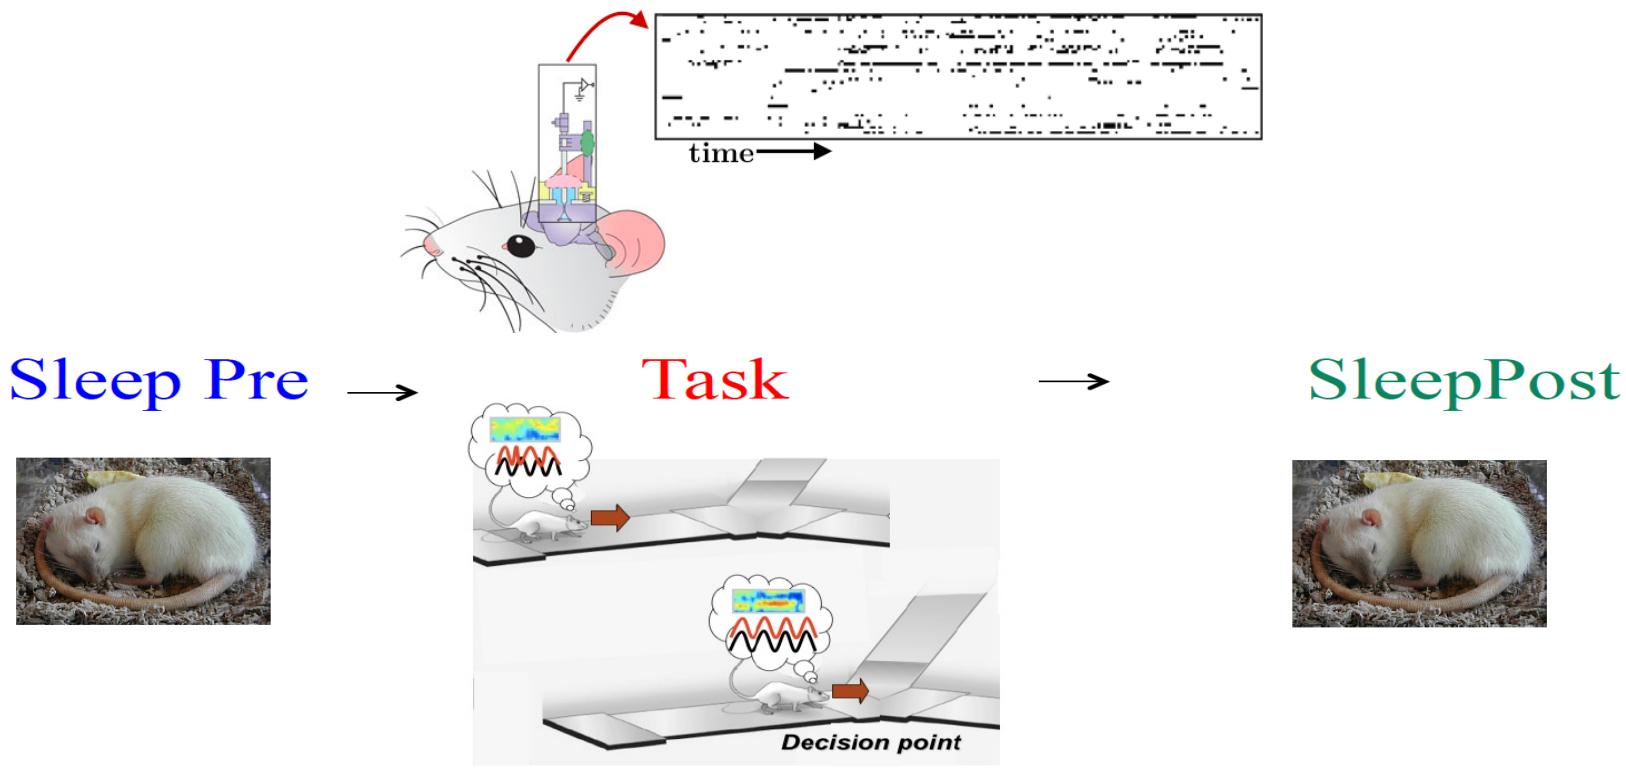
Tutorials based on programming in **Python** to **analyse** real data:

Data: Neural Recordings, Images data, behavioral recordings

You need a laptop with jupyter notebook installed (eg. by anaconda): For tomorrow:

See: <https://moodle.psl.eu/enrol/index.php?id=34555>

Tutorial 4: Principal Component Analysis to study memory consolidation during the Sleep after a task



Replay of rule-learning related neural patterns in the prefrontal cortex during sleep (2009)

*nature
neuroscience*

Adrien Peyrache¹, Mehdi Khamassi^{1,2}, Karim Benchenane¹, Sidney I Wiener¹ & Francesco P Battaglia^{1,3}

How to follow the Course

Moodle contains:

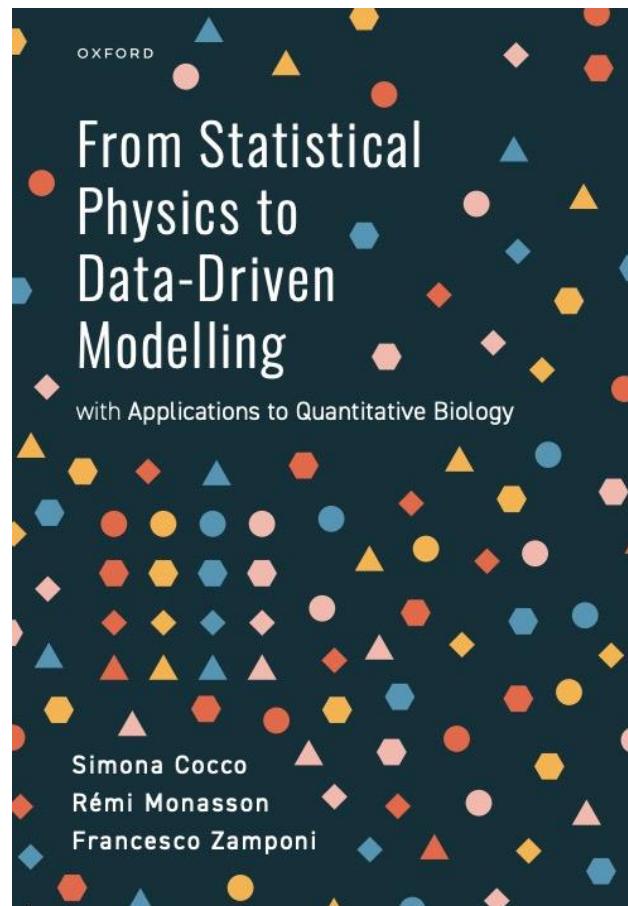
<https://moodle.psl.eu/enrol/index.php?id=34555>

- A pdf version of the book
- Each week: The slides,
- The tutorials (data and starting notebook) ,
- The final notebook after the tutorials
- Check you have access to it.

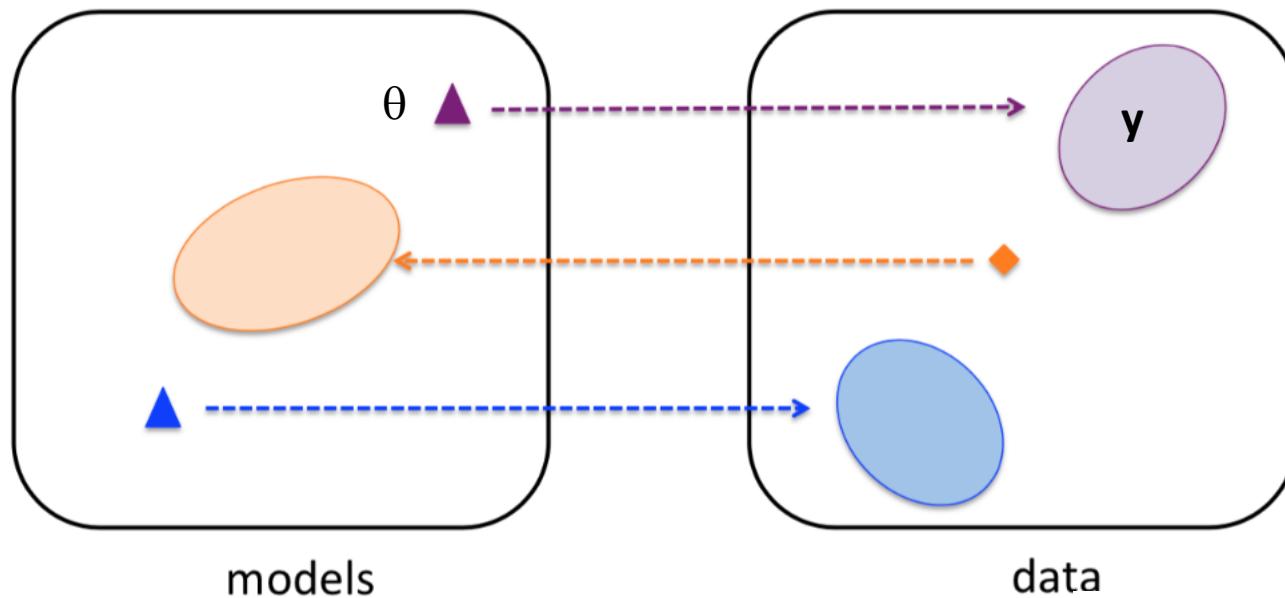
Class each Wednesday 8h30-10h30

Tutorials each Thursday 8h30-10h30

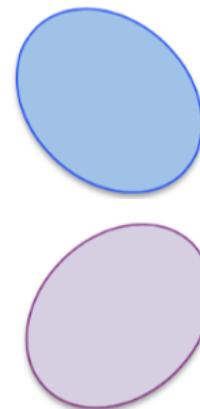
- Install Jupyter Notebook (Anaconda) Python 3
- Revise lectures before the tutorials
- Attendance is mandatory.
- Final scores: One homework before christmas,
 - + final examination : python program based on the tutorial in class and a session of theoretical questions



Bayesian Inference



- Each model \triangle defines a distribution of possible data
- Each data \diamond defines a distribution of possible models



Probabilities and conditional probabilities

Dice : faces $y = 1, 2, 3, 4, 5, 6$ (here, $y = 3$)



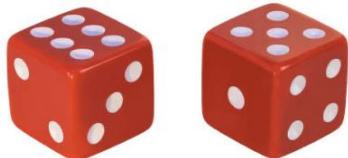
Notation: Probability (face = y) $\equiv p(y)$

Properties : $p(y) \geq 0$, $\sum_{y=1,\dots,6} p(y) = 1$

Unbiased dice : each face is equally likely through a draw

$$p(1) = p(2) = \dots = p(6) = \frac{1}{6}$$

Probabilities and conditional probabilities



Two dices : faces y_1 and y_2 (here, $y_1 = 6$, $y_2 = 5$)

Joint probability (1st dice = y_1 & 2nd dice = y_2) $\equiv p(y_1, y_2)$

Two important considerations:

- Care about one event only -> definition of **marginal probability**

$$p_1(y_1) = \sum_{y_2} p(y_1, y_2) , \quad p_2(y_2) = \sum_{y_1} p(y_1, y_2)$$

- Ask whether events are **independent**, i.e. are the two dices correlated or not?

$$p(y_1, y_2) = p_1(y_1) \times p_2(y_2) ?$$

An example of dependent events

Box A

20 plain cookies
+
20 chocolate cookies

Box B

10 plain cookies
+
30 chocolate cookies

One picks up uniformly at random one box and one cookie out of this box

Box: $x_1 = A \text{ or } B$ $p_1(A) = p_1(B) = \frac{1}{2}$

Cookie type: $x_2 = \text{plain or chocolate}$ $p_2(\text{plain}) = \frac{3}{8}$
 $p_2(\text{chocolate}) = \frac{5}{8}$

Conditional Probability

but $p_2(\text{chocolate} | A) = \frac{1}{2} < p_2(\text{chocolate}) = \frac{5}{8} < p_2(\text{chocolate} | B) = \frac{3}{4}$

An example of dependent events

Box A

20 plain cookies
+
20 chocolate cookies

Box B

10 plain cookies
+
30 chocolate cookies

One picks up uniformly at random one box and one cookie out of this box

Box: $\theta = A \text{ or } B$ $p_1(A) = p_1(B) = \frac{1}{2}$

Cookie type: $y = \text{plain} \text{ or } \text{chocolate}$ $p_2(\text{plain}) = \frac{3}{8}$
 $p_2(\text{chocolate}) = \frac{5}{8}$

- When the events are **dependent**:

$$p(y, \theta) = p_1(\theta) \times p_2(y | \theta)$$

Characterization of dependent events

Definition of **conditional probability**:

$$p_2(y|\theta) \equiv \frac{p(y, \theta)}{p_1(\theta)}$$

Exercise 1: check that marginal probability over the boxes give the probabilities of Chocolate of plain cookies computed in previous slide .

Exercise 2: check that conditional probability is normalized.

Box A

20 plain cookies
+
20 chocolate cookies

$$p_2(y=\text{plain} | \theta=A) = 1/2,$$
$$p_2(y=\text{chocolate} | \theta=A) = 1/2$$

Box B

10 plain cookies
+
30 chocolate cookies

$$p_2(y=\text{plain} | \theta=B) = 1/4,$$
$$p_2(y=\text{chocolate} | \theta=B) = 3/4$$

Characterization of dependent events

Definition of **conditional probability**:

$$p_1(\theta|y) \equiv \frac{p(y, \theta)}{p_2(y)} \equiv \frac{p_2(y|\theta) \times p_1(\theta)}{p_2(y)}$$

*Probability that the Box is A or B given that
I have drawn a plain cookie*

Box A

20 plain cookies
+
20 chocolate cookies

Box B

10 plain cookies
+
30 chocolate cookies

$$p_2(\theta=A | \text{plain}) = ?$$

$$p_2(\theta=B | \text{plain}) = ?$$

Characterization of dependent events

Definition of **conditional probability**:

$$p_1(\theta|y) \equiv \frac{p_2(y|\theta) \times p_1(\theta)}{p_2(y)}$$

1/2
3/8

Probability that the Box is A or B given that I have drawn a plain cookie

Box A

20 plain cookies
+
20 chocolate cookies

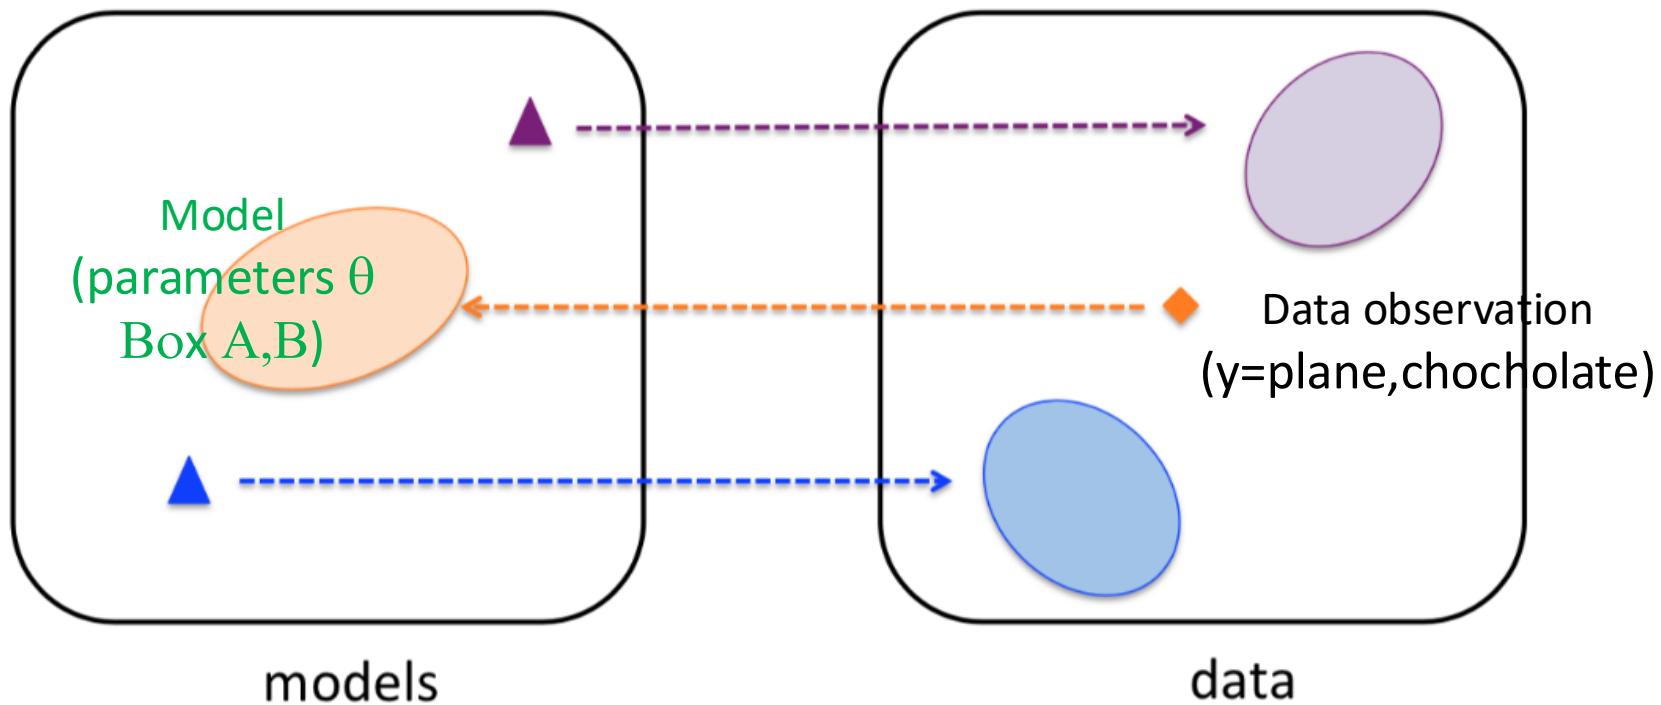
Box B

10 plain cookies
+
30 chocolate cookies

$$p_2(\theta=A | \text{plain}) = \frac{\frac{1}{2} \times \frac{1}{2}}{\frac{3}{8}} = \frac{2}{3}$$

$$p_2(\theta=B | \text{plain}) = \frac{\frac{1}{4} \times \frac{1}{2}}{\frac{3}{8}} = \frac{1}{3}$$

Bayesian Inference



$$p(\theta|y) = \frac{p(y|\theta) \times p(\theta)}{p(y)}$$

Likelihood of data given the model

A priori distribution of the model parameters

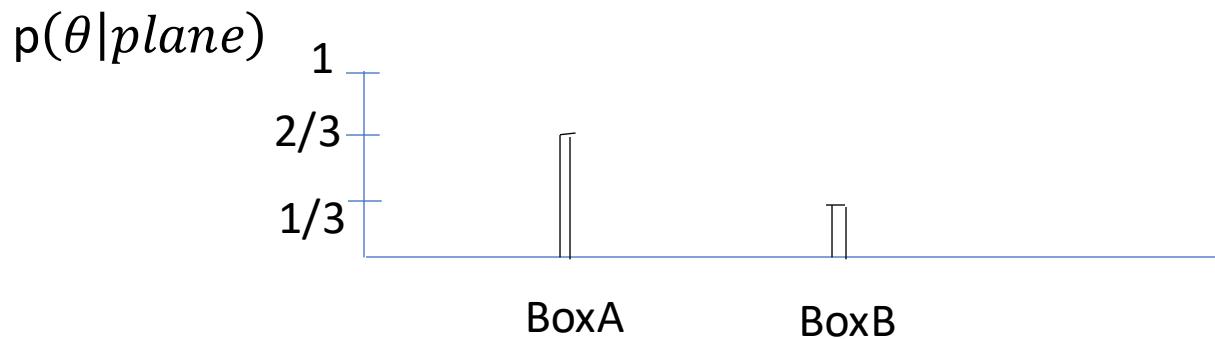
Posterior of model parameter given the data

Probability of that data: Evidence

An Illustration: Model (Box) Inference from the Bayes formula

Suppose you have drawn a plane cookie

Posterior of the model parameters given the data



- Most Probable model: Box A $p_2(\text{plane})=1/2$, $p_2(\text{pchocolate})=1/2$

From the model you can predict the probability of having a chocolate cookie if you draw again from the same box

- Probability that the prediction is wrong: $1/3$

How you can increase the confidence in your model?

->lower the probability of a wrong prediction

Laplace and the birth rate of boys & girls

Historical example: « proof » by Laplace that the female and male birth rates are different

Data: Nbs of girls born in Paris from 1745 to 1770 : 245,945
... boys ... : 251,527

y = nb. of female births, M = total number of births

Laplace and the birth rate of boys & girls

Historical example: « proof » by Laplace that the female and male birth rates are different

Data: Nbs of girls born in Paris from 1745 to 1770 : 245,945
... boys ... : 251,527

y = nb. of female births, M = total number of births

Inference: θ = probability that a newborn baby is a girl

- Prior distribution: uniform over θ in $[0;1]$

INTUITION: $\theta \sim$

Laplace and the birth rate of boys & girls

Historical example: « proof » by Laplace that the female and male birth rates are different

Data: Nbs of girls born in Paris from 1745 to 1770 : 241,945
... boys ... : 251,527

y = nb. of female births, M = total number of births

Inference: θ = probability that a newborn baby is a girl

- Prior distribution: uniform over θ in [0;1]

INTUITION: $\theta \sim \frac{241,945}{493,472} = 0.4903$

Laplace and the birth rate of boys & girls

Historical example: « proof » by Laplace that the female and male birth rates are different

Data: Nbs of girls born in Paris from 1745 to 1770 : 241,945
... boys ... : 251,527

σ = nb. of female births, n = total number of births

Inference: θ = probability that a newborn baby is a girl

- Prior distribution: uniform over θ in $[0;1]$
- Likelihood: $p(y|\vartheta) = \binom{M}{y} \vartheta^y (1 - \vartheta)^{M-y}$

Binomial Distribution.

Laplace and the birth rate of boys & girls

Historical example: « proof » by Laplace that the female and male birth rates are different

Data: Nbs of girls born in Paris from 1745 to 1770 : 241,945
... boys ... : 251,527

σ = nb. of female births, n = total number of births

Inference: θ = probability that a newborn baby is a girl

- Prior distribution: uniform over θ in $[0;1]$

- Likelihood: $p(y|\theta) = \binom{M}{y} \vartheta^y (1 - \vartheta)^{M-y}$

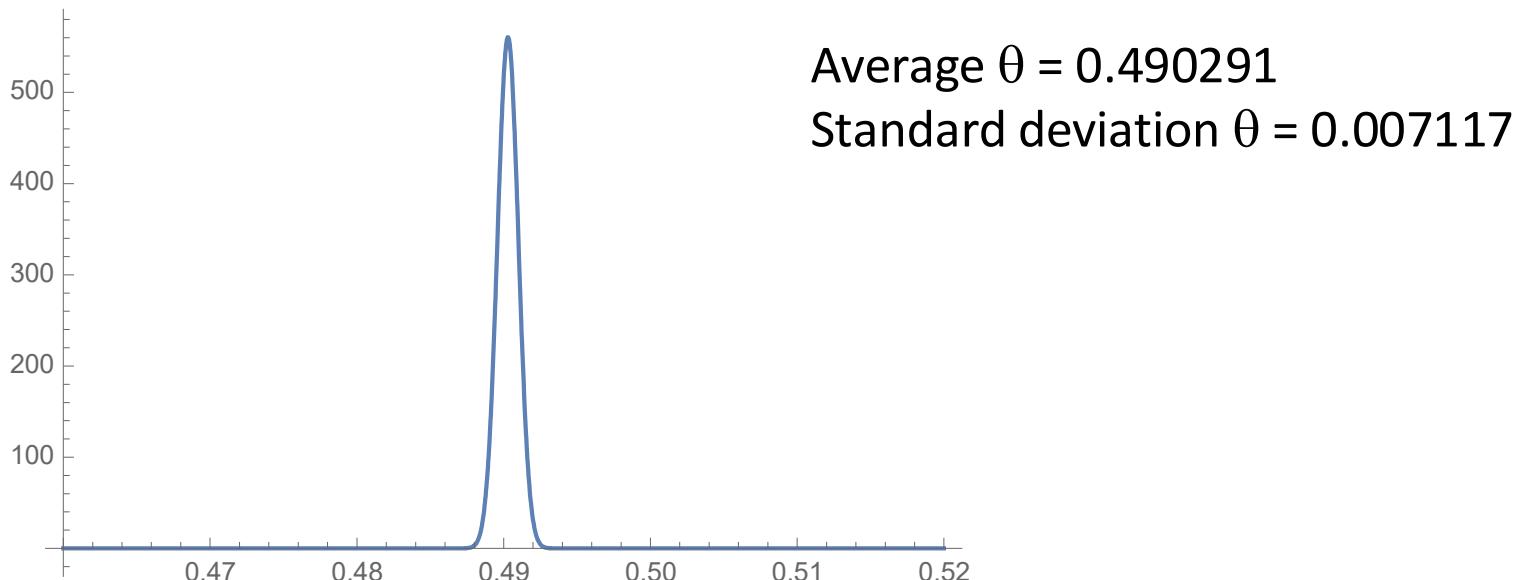
- Bayes:
$$p(\theta|y) = \frac{p(y|\theta) \times p(\theta)}{p(y)}$$

Uniform in interval $[0,1]$

\leftarrow Cst $\int_0^1 d\theta \vartheta^y (1 - \vartheta)^{M-y}$

Laplace and the birth rate of boys & girls

Posterior distribution:



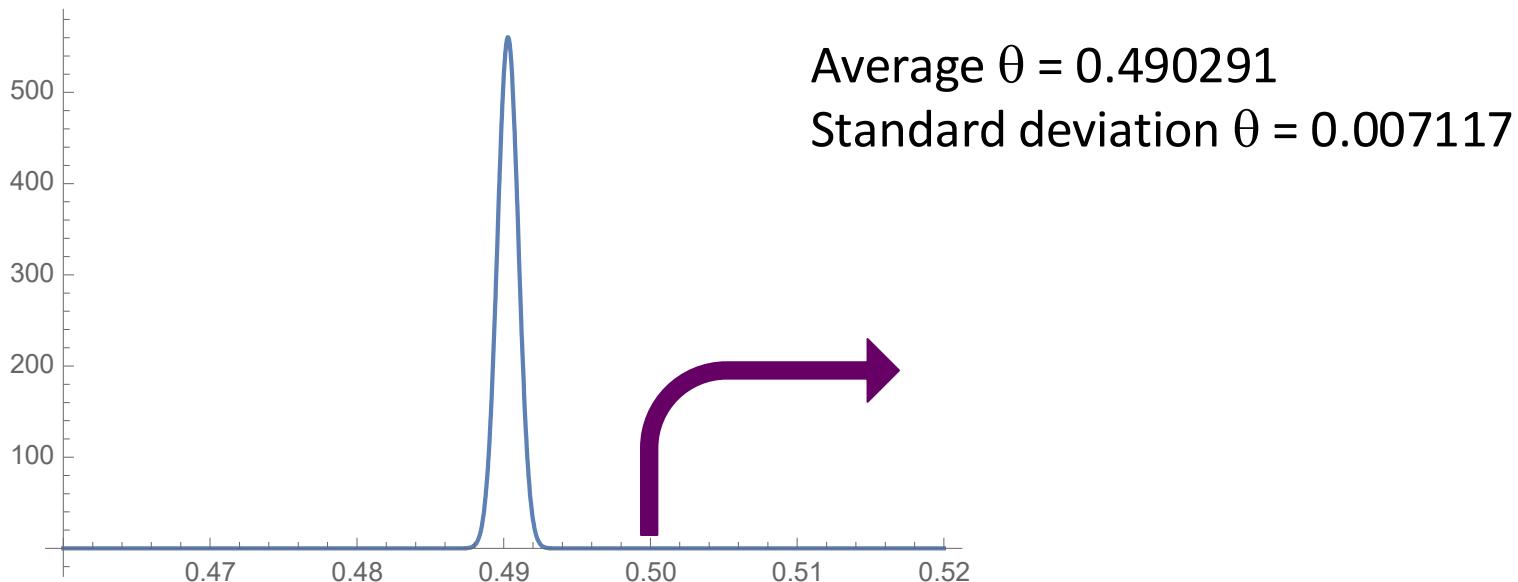
- Bayes: $p(\theta|y) = \frac{p(y|\theta) \times p(\theta)}{p(y)}$

Uniform in interval $[0,1]$

Cst $\int_0^1 d\theta \theta^y (1 - \theta)^{M-y}$

Laplace and the birth rate of boys & girls

Posterior distribution:



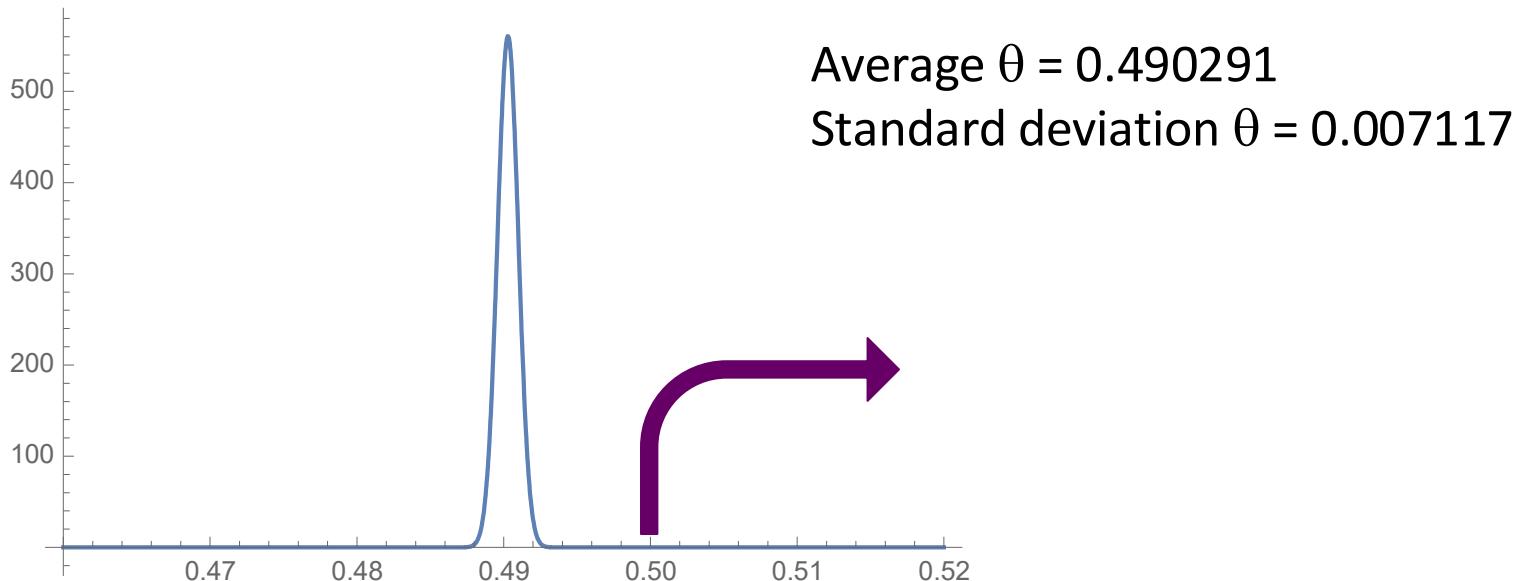
Probability that θ is equal or larger than 0.5 =

$$\int_{0.5}^1 d\theta \ p(\theta|y) \approx 10^{-42}$$

Extremely unlikely!

Laplace and the birth rate of boys & girls

Posterior distribution:



$$\text{Probability that } \theta \text{ exceeds } 0.5 = \int_{0.5}^1 d\theta \ p(\theta|y) \approx 10^{-42}$$

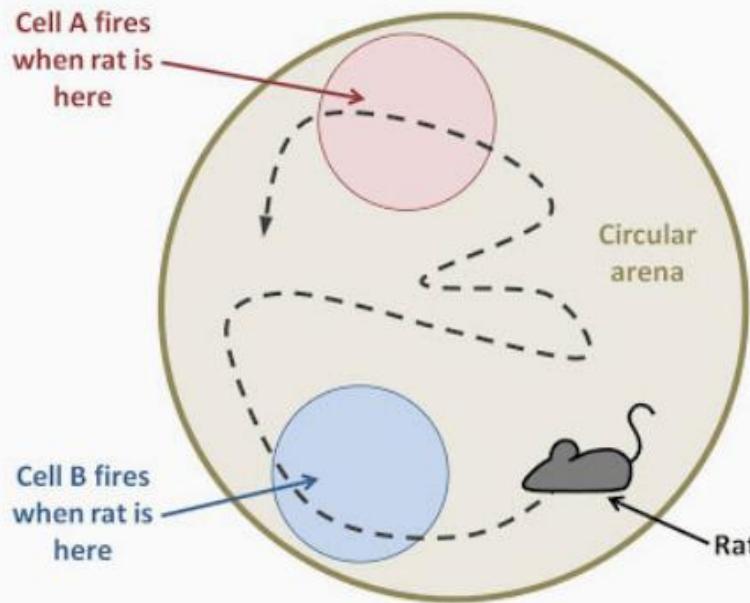
Extremely unlikely!

- In the tutorial tomorrow you will compute the posterior distribution, the most likely value, and average value for the spiking rate of a neuron from a neural recording (Poissonian Distribution)

Application:
**Decoding of position from recorded activity in
the Hippocampus (HPC)**

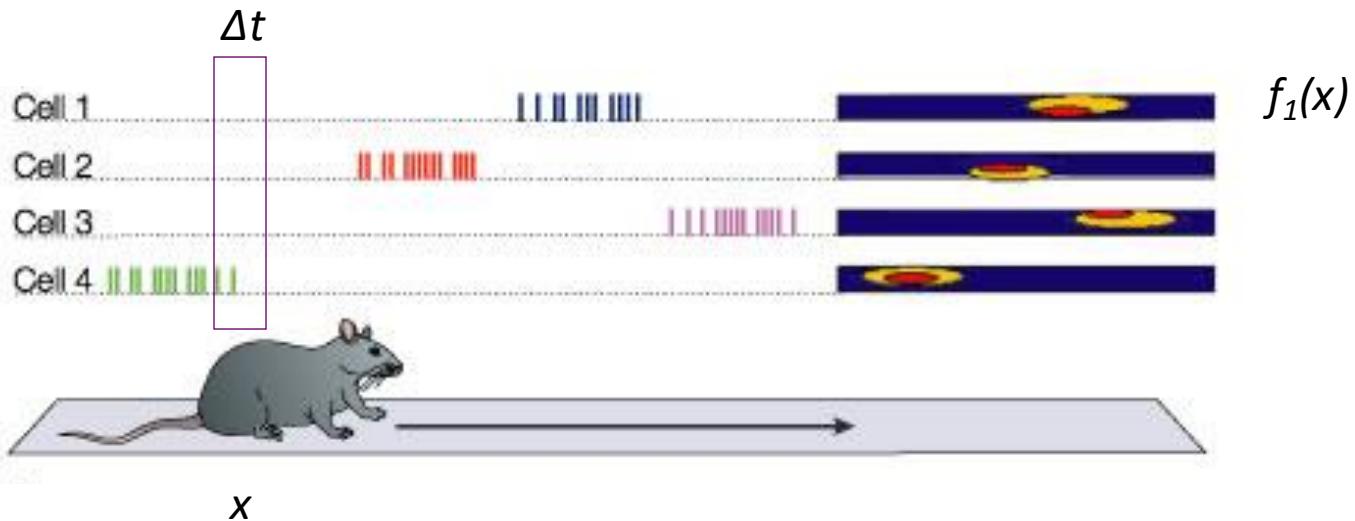
Decoding of position from HPC activity

Place cells in the hippocampus regions CA1 and CA3 present spatially-located firing fields.



- Place fields are retrieved when the animal is placed in the same environment after days
- Stable in dark and against limited changes of environment

Decoding of position from HPC activity



- Assume place cells are Poisson :

Probability of y spikes for cell i :

$$p_i(y_i|x) = e^{-f_i(x)\Delta t} \frac{(f_i(x) \Delta t)^{y_i}}{y_i!}$$

- First one infers the spiking rates of the cells at each position $f_i(x)$

- Likelihood for the population activity:

Assume place cell activities are independent (conditional to position x):

$$P(\{y_i\}|x) = \prod_i p_i(y_i|x)$$

Decoding of position from HPC activity

- Prior over trajectory: $P_{prior}(x_t)$ = uniform over the environment

- Find most likely position:

$$MAP=ML\ decoding \quad x_t = \operatorname{argmax}_x [P(\{y_i\}|x) \times P_{prior}(x_t)]$$

Decoding of position from HPC activity

- Prior over trajectory: $P_{prior}(x_t)$ = uniform over the environment
- Find most likely position:

$$MAP=ML\ decoding \quad x_t = \operatorname{argmax}_x [P(\{y_i\}|x) \times P_{prior}(x_t)]$$

Example: CA1 recording in freely moving rats, 60 cm squared box, 34 cells
Place field: 400 squared bins per cell (discretized x – in 2D)
(data from K. Jezek)

Average error ($Dt = 120$ ms): 12 cm ...

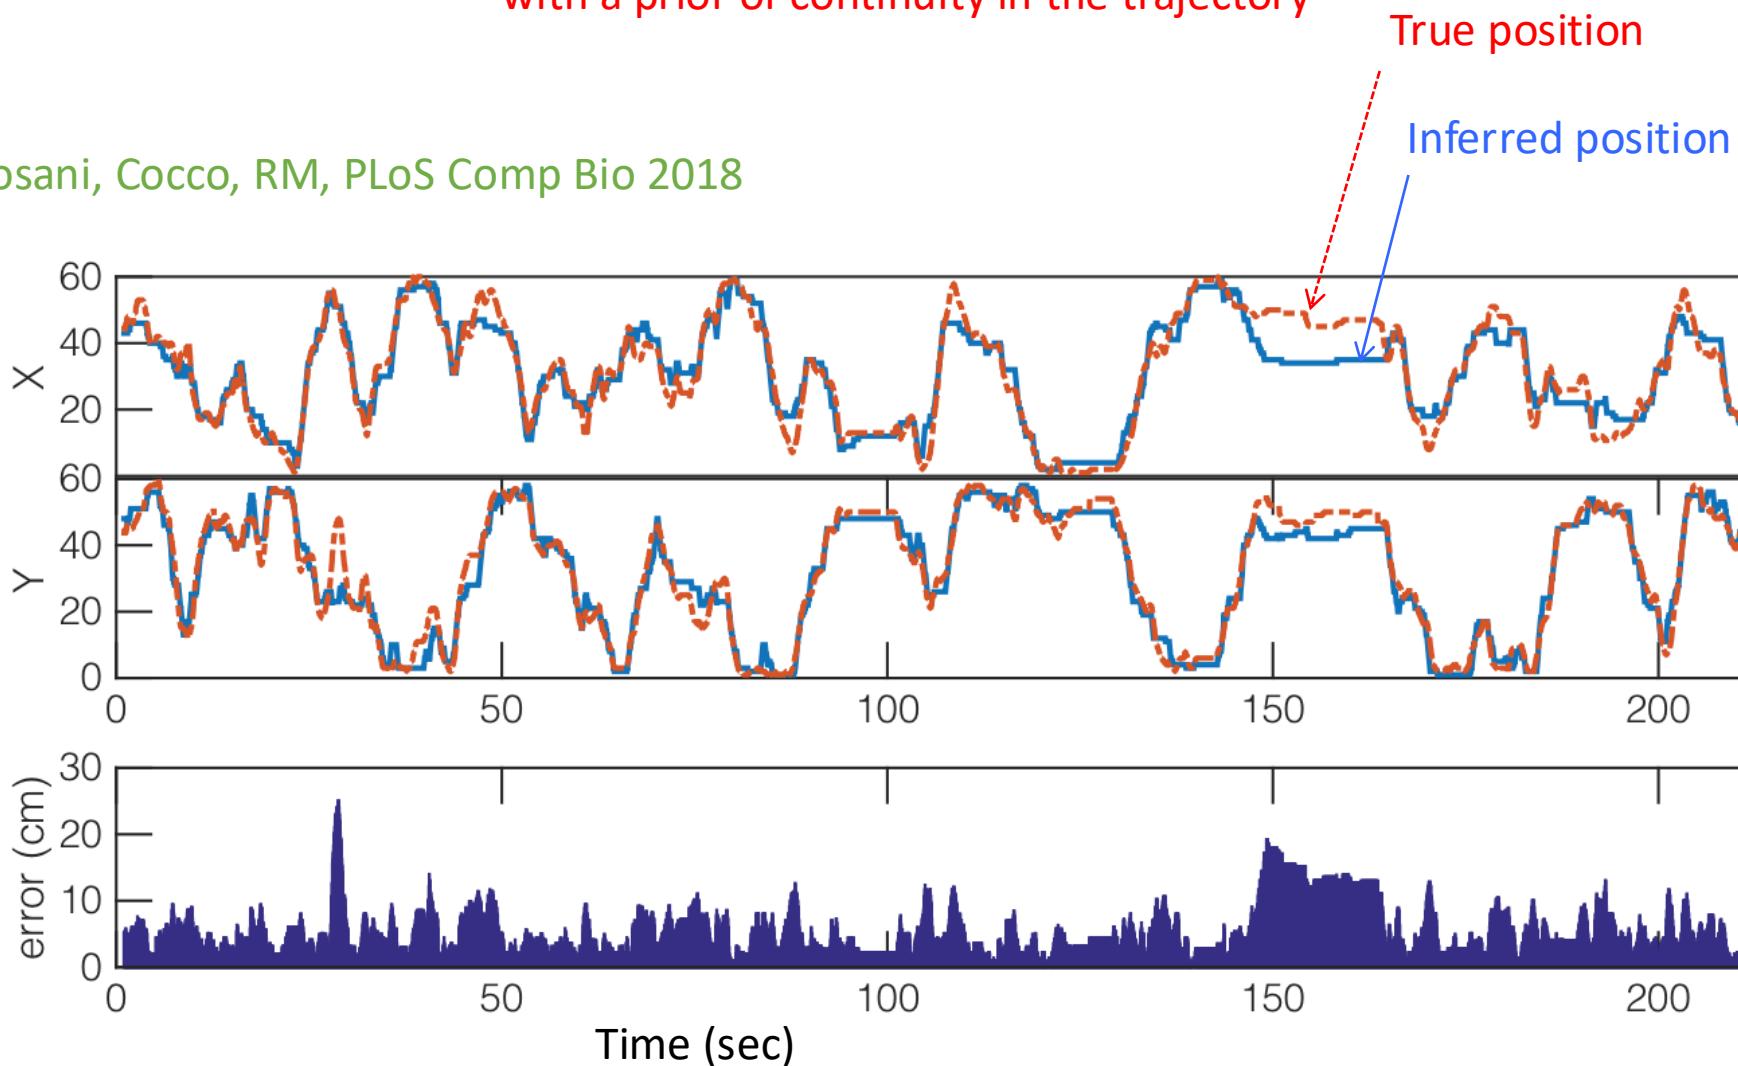
- Some regions in space are barely covered by place fields ...
- Synthetic data study: error on positional decoding decreases as $(nb.\ cells)^{-1/2}$

How to improve precision?

Decoding of position from HPC activity

with a prior of continuity in the trajectory

Posani, Cocco, RM, PLoS Comp Bio 2018



Average error ($Dt = 120$ ms): 5.5 cm (compared to 12 cm)

Hippocampal Reactivation of Random Trajectories Resembling Brownian Diffusion

Authors

Federico Stella, Peter Baracskay,
Joseph O'Neill, Jozsef Csicsvari

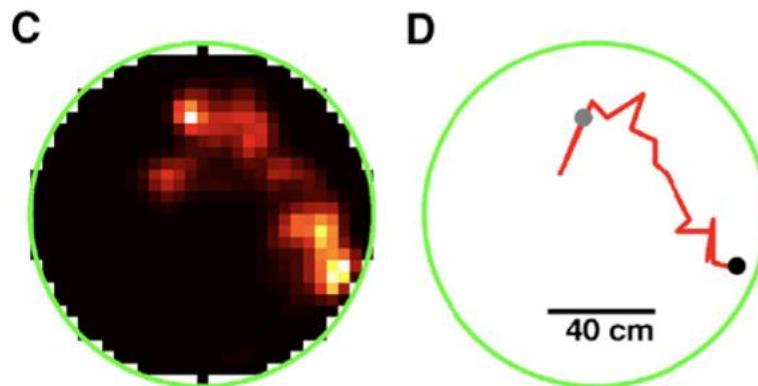
Activity During Sleep:



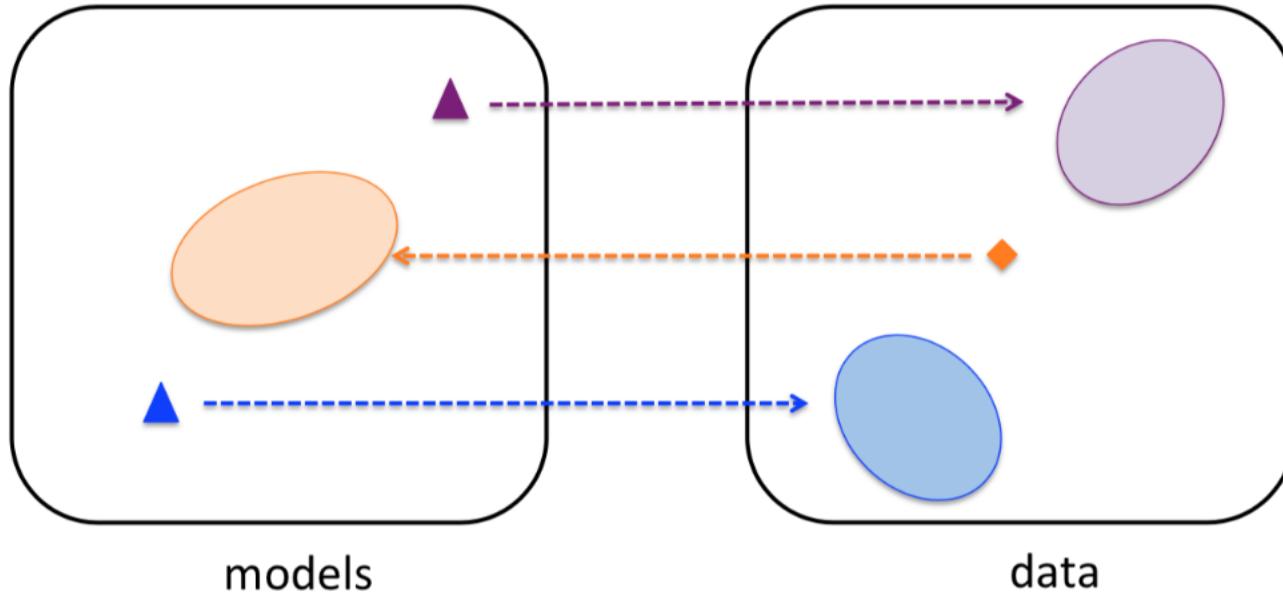
Decoding
model form
recordings
during awake
exploration

In Brief

Stella et al. examine the dynamical properties of reactivated spatial trajectories in the hippocampus following non-stereotypical exploration and find that reactivated trajectories are governed by a Brownian diffusion process that occur at varying lengths and timescales, without directly reflecting behavior.



Take home Message

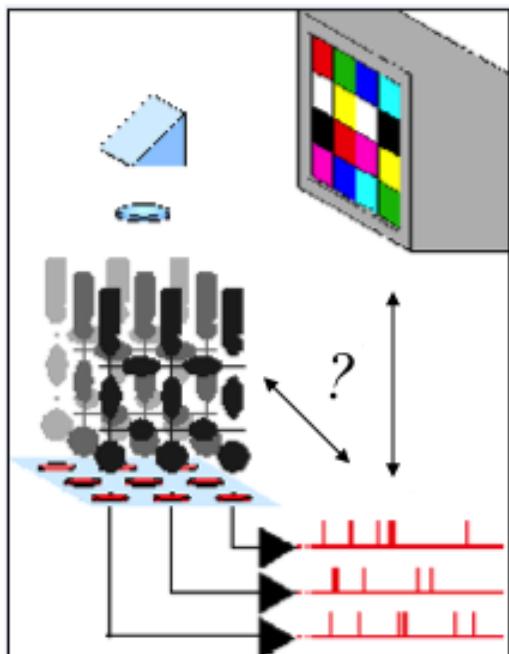


- Use Bayes law to derive the posterior distribution of parameter value given the observable.
Key concepts: Posterior probability of a model given the data, Likelihood of data given a model, Prior probability of model parameters,

Tutorial: A Poisson process to model the spiking activity of a neuron and inference of the spiking frequency



In vitro recording on retina



tutorial:
40 ganglions, salamandra retina
natural movies. Recordings last for
about 1 hour

Vol 440 | 20 April 2006 | doi:10.1038/nature04701 nature

ARTICLES

Weak pairwise correlations imply strongly correlated network states in a neural population

Elad Schneidman^{1,2,3}, Michael J. Berry II², Ronen Segev² & William Bialek^{1,3}

Tutorial

Modeling Spike trains with a Poissonian process



$$p(y|\theta) = \frac{\theta^y e^{-\theta}}{y!} \quad \Theta=f \Delta t$$

- Plot the Posterior distribution, Infer from a real spike train $\theta^* \langle \theta \rangle \sigma_\theta^2$
- For different recording times T .
- Is the Poisson process a good model for the data?

Before the tutorial of tomorrow:

- Install Jupyter Notebook (Anaconda) Python 3
- Study slides of the class today, jupyter notebook of the class
- Book's Chapter 1

Supplementary Slides

5 Important Probability Distributions For Random variables

- Uniform Distribution

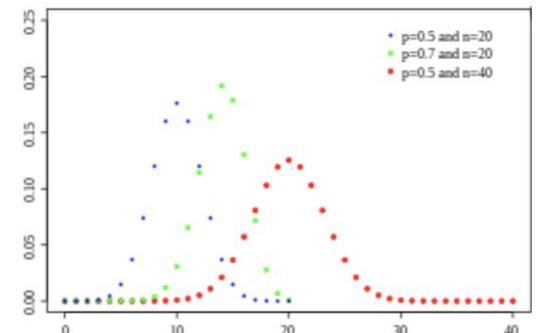
$$y \in \{a_1, a_2, \dots, a_q\} \quad p(y) = 1/q$$



- Binomial Distribution

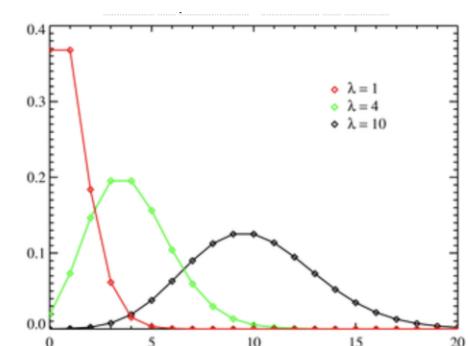
$$p(y|\theta) = \binom{M}{y} \theta^y (1-\theta)^{M-y}$$

$$y \in \{0, 1, 2, M\} \in \mathbb{N}$$



- Poisson Distribution

$$p(y|\lambda) = \frac{\lambda^y e^{-\lambda}}{y!} \quad y \in \{0, 1, 2, 3, \dots, \infty\} \in \mathbb{N}$$

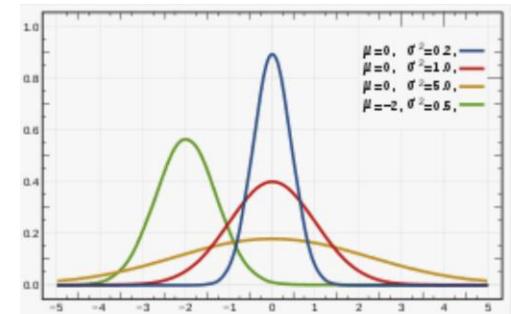


Note, the probability is always $[0,1]$ and normalized to 1 !!

5 Important Probability Distributions For Random variables

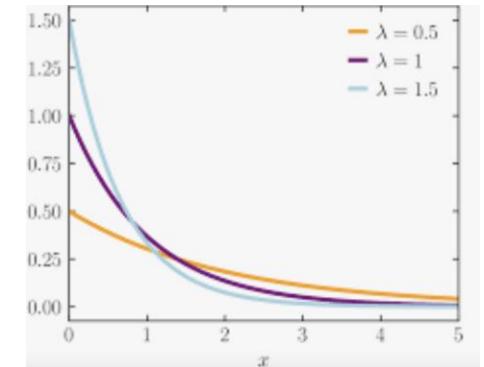
- Normal/Gaussian Distribution

$$y \in \mathbb{R} \quad p(y|\mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{(y-\mu)}{\sigma}\right)^2}$$



- Exponential Distribution

$$y \in \mathbb{R}, \text{ positives} \quad p(y|\lambda) = \lambda e^{-\lambda y}$$



Note, the probability is always [0,1] and normalized to 1 !!

Bayes Theorem in Inference

Suppose we have an ensemble of M data points $\mathbf{y}_i \in \mathbb{R}^L$:

Generated with a model with \mathbf{D} unknown parameters $\boldsymbol{\theta} \in \mathbb{R}^D$

Bayes' rule:

$$p(\theta|Y) = \frac{p(Y|\theta)p(\theta)}{p(Y)}.$$

The Evidence ensures the normalization of the Posterior

$$p(Y) = \int d\boldsymbol{\theta} p(Y|\boldsymbol{\theta})p(\boldsymbol{\theta}) .$$

Technical Tricks:

Use the property of the Beta Distribution

$$\text{Beta}(\theta; \alpha, \beta) = \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}}{B(\alpha, \beta)}$$

$$B(\alpha, \beta) = \int_0^1 d\theta \, \theta^{\alpha-1} (1-\theta)^{\beta-1} = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}$$

In our case: $\alpha = y + 1$ $\beta = M - y + 1$,

For α integer

$$\Gamma(\alpha) = \int_0^\infty d\theta \, \theta^{\alpha-1} e^{-\theta} = (\alpha-1)!$$